

TRANSCRIBE PROJECT DOCUMENTS

Implementation Plan

Document 06

Status	Draft v1
Team	1 developer/reviewer. CPU laptop + Kaggle GPU
Related	All previous docs. Each task ID maps to PRD/TRD requirement IDs where applicable

Contents

1	Phase 0: Foundations (≈ 5 days)	1
2	Phase 1: Alignment pipeline (≈ 7 days)	2
3	Phase 2: Review UI v2 (≈ 6 days)	2
4	Phase 3: Datasets + zero-shot baselines (≈ 5 days)	3
5	Phase 4: Fine-tune v1 on Kaggle (≈ 6 days, GPU-quota bound)	4
6	Phase 5: Transcribe service (≈ 8 days; can overlap with P4)	4
7	Phase 6: Feedback loop + model v2 (≈ 10 days)	5
8	Definition of done (every task)	5
9	Experiment log template (docs/experiments/EXX_name.md)	5
10	Risks to the schedule	5
11	Immediate next steps	6

The phases are ordered so that each one produces something usable and measurable. Estimates are working days for one person. Review hours for labelling are listed separately because they run in parallel with coding.

1 Phase 0: Foundations (≈ 5 days)

Goal: a clean base that stops the known bugs from spreading into new work.

ID	Task	Acceptance
P0.1	Make transcribe an installable package: drop the sys.path hack and the src. imports; entry points transcribe-api, transcribe-worker, transcribe-cli	pip install -e . then python -c "import transcribe" works; all scripts run via python -m transcribe...
P0.2	Move models/, training/, data/, script/{train,evaluate,infer}.py into transcribe/ legacy/ unchanged	Legacy train still runs on a toy dataset

ID	Task	Acceptance
P0.3	Fix: <code>transcribe_current</code> must not overwrite <code>segment["transcript"]</code> ; store the hypothesis only	Test: after Transcribe, the transcript equals the previous value
P0.4	Fix: grouped split by recording in the legacy <code>prepare_dataset</code> ; build the vocabulary from train only	Test: no recording in two splits
P0.5	Add <code>pytest</code> , <code>ruff</code> , and a GitHub Actions workflow (Ubuntu, Py 3.10)	CI green on the branch
P0.6	SQLAlchemy models + Alembic initial migration from 05 Backend Schema	<code>alembic upgrade head</code> creates all tables; invariant tests 1, 2, 8 pass
P0.7	<code>transcribe.db.migrate_json</code> imports the existing <code>data/projects/*</code> manifests	Approved segments are preserved; candidates are dropped with a count printed
P0.8	Extend <code>docs/STYLE_GUIDE.md</code> (v0.1 drafted) (verbatim rules, hybrid words, numbers, markers, language tags) with 20 worked examples from real transcripts	Reviewed by you; referenced from the UI help
P0.9	Split requirements into <code>requirements/*.txt</code> as in the TRD §7	Fresh <code>venv</code> installs <code>core+corpus+ui</code> on Windows and Linux

2 Phase 1: Alignment pipeline (≈7 days)

Goal: DOCX + audio in, pre-filled segments out.

ID	Task	Acceptance
P1.1	<code>corpus/ingest.py</code> : <code>sha256</code> , <code>ffmpeg</code> normalise to 16 kHz mono, store the original	Re-importing the same file → 409
P1.2	Move the parsers into <code>corpus/parsers/</code> ; add <code>text_align</code> + <code>word_map</code> generation (<code>corpus/textnorm.py</code>)	Unit tests: annotations removed, number verbalisation (EN via <code>num2words</code> + <code>numbers_sw.py</code> , candidates scored by the aligner, <code>number_check</code> flag on ties), word map round-trips to the raw text
P1.3	<code>corpus/align.py</code> : an Aligner protocol + an MMS implementation (windowed emissions, turn-anchored alignment, per-word scores)	60 s fixture aligns with mean word boundary error < 100 ms vs. hand labels; a 1 h recording aligns on CPU in < 20 min
P1.4	<code>corpus/segment.py</code> : the TRD §3 policy (≤ 30 s, 1.5 s minimum, $\text{gap} \geq 300$ ms, turn-aware, padding)	Property tests: no segment > 30 s, no word cut, full coverage of aligned words
P1.5	Auto-flags (<code>low_alignment</code> , <code>unclear_marker</code> , <code>overlap_marker</code>) + <code>priority.py</code>	Unit tests on synthetic inputs
P1.6	jobs worker loop with claim, heartbeat and retry; job kinds <code>parse_transcript</code> , <code>align</code>	Killing the worker mid-job → the job is re-queued and completes on restart
P1.7	Calibrate 0: align one real recording, hand-check 50 segments, and set the <code>low_alignment</code> threshold to catch $\geq 90\%$ of bad segments	The threshold and its precision/recall are recorded in <code>docs/experiments/E1_alignment.md</code>

3 Phase 2: Review UI v2 (≈6 days)

Goal: reviewing becomes verification, driven from the keyboard.




Figure 1: Indicative schedule for one developer (working days).

ID	Task	Acceptance
P2.1	FastAPI app + the /recordings, /segments, /review/ queue routes; revision endpoint with 409 on conflict	API tests including the conflict case
P2.2	Gradio shell with 3 tabs and theme tokens from 03 UI/UX §2, mounted in FastAPI	Loads at 127.0.0.1:8000
P2.3	Import wizard + recording overview with confidence timeline	Import → align → overview with no CLI
P2.4	Segment reviewer: pre-filled text, flags, boundary nudging with slice re-render, approve & next, history/restore	A reviewer can complete 20 segments using only the keyboard
P2.5	Hotkey JS layer	All keys in the §4.5 table work
P2.6	Language span tagging (P1 priority; may slip to P6)	Spans persist per revision
P2.7	Measure throughput on one real recording: minutes of review per audio minute vs. the old type-from-scratch UI	Recorded in docs/experiments/E2_review_speed.md; target ≥ 3× faster

Labelling track (parallel, starts at the end of P2): review enough recordings to build **test_v1** ≈ 3 h (≥ 15 recordings, stratified), **dev_v1** ≈ 2–3 h (≥ 10 recordings), and **train** ≥ 20 h for v1. Review the test/dev recordings first; they matter most and must be the most carefully checked.

4 Phase 3: Datasets + zero-shot baselines (≈5 days)

Goal: know where we stand before training anything. This is the most instructive phase for learning the process.

ID	Task	Acceptance
P3.1	datasets/build.py: filter spec, grouped split, frozen test_v1, sha256, invariants 4–7	Tests for leakage and consent enforcement
P3.2	datasets/export.py: audiofolder export + dataset card	load_dataset("audiofolder", ...) loads it on Kaggle
P3.3	evaluation/normalizer.py + metrics.py (WER, CER, per-language, switch-point, hallucination) with hand-computed tests	100% test coverage on metric functions

ID	Task	Acceptance
P3.4	inference/engine.py (faster-whisper) + register the stock models as zero-shot baselines	whisper-small, whisper-medium and large-v3-turbo are registered
P3.5	Experiment E3 (zero-shot): evaluate each stock model on dev with language sw, en and auto	A table in docs/experiments/E3_zero_shot.md with an error analysis of 30 worst segments (translation? hallucination? spelling of hybrids?)

5 Phase 4: Fine-tune v1 on Kaggle (≈6 days, GPU-quota bound)

ID	Task	Acceptance
P4.1	finetune/whisper_train.py: HF Seq2SeqTrainer, config YAML, fixed < sw > token, SpecAugment, non-speech samples, public-data mixture with weights, resume support, run_manifest.json	Dry run on 50 samples on CPU completes 5 steps
P4.2	Kaggle notebook notebooks/kaggle/finetune_whisper.ipynb (thin wrapper) + a private dataset upload of the source package	Runs end to end on a T4
P4.3	Experiment E4: whisper-small full fine-tune, in-domain only vs. in-domain + public mixture	Dev WER for both; the winner is kept
P4.4	Experiment E5: language token sw vs en vs none (on the E4 winner, 1.5k steps each)	Decision recorded
P4.5	Register + CT2 int8 convert + evaluate on test_v1	whisper-small-kct-v1 has a test report; gate: clean test WER ≤ 20% or a documented plan for what to change
P4.6	(Learning track, optional) finetune/w2vbert_ctc_train.py + a KenLM 4-gram trained on approved train text	Comparison row in the evaluation table

6 Phase 5: Transcribe service (≈8 days; can overlap with P4)

Goal: the v1 product surface.

ID	Task	Acceptance
P5.1	inference/vad.py (Silero), longform.py (windowing ≤ 30 s), guards.py (no-speech, compression ratio, n-gram loop)	A 10 min music-only file produces ≤ 10 words
P5.2	transcribe job with incremental segment writes and progress/ETA	A 1 h file on a 4-core CPU finishes with RTF ≤ 0.5
P5.3	Transcribe tab: upload, jobs list, editor with autosave, low-confidence highlighting	Manual walkthrough of a 20 min file
P5.4	Exports TXT/SRT/VTT/JSON (+ DOCX P1); edited text wins over raw output	Golden-file tests for each format
P5.5	Model registry screen, promotion gate, hot reload	Promoting v1 switches Transcribe without a restart
P5.6	Packaging: Dockerfile (CPU), make run, a first-run setup checklist (ffmpeg, model weights)	Fresh machine to a working app in < 15 min following the README

v1 release = end of P5. A working prototype: upload, transcribe, edit and export, with a fine-tuned model

and a test report.

7 Phase 6: Feedback loop + model v2 (≈10 days)

ID	Task	Acceptance
P6.1	“Promote to corpus” from Transcribe (consent-checked, edited segments only)	Invariant tests pass
P6.2	VAD-only import for untranscribed recordings + model pre-fill	Flow B works end to end
P6.3	Active-learning priority from model uncertainty; weekly batch job	The queue orders by uncertainty
P6.4	Domain eval sets: ≥ 1 h each of TV and church audio, reviewed, stored as external test sets	Per-domain WER in reports
P6.5	Experiment E6: whisper-large-v3-turbo with LoRA (r=32, attention + FFN) vs. v1	Promoted only if it passes the gate; RTF is checked on CPU
P6.6	Learning curve plot: test WER vs. training hours across snapshots	docs/experiments/E6_learning_curve.md
P6.7	Optional: pyannote diarization for Speaker 1/2 labels in the Transcribe output	Behind a feature flag

8 Definition of done (every task)

- Code under src/transcribe/... with type hints. Unit tests are added, and CI is green.
- No change to reviewer text by any automated path (invariant 3 test stays green).
- Any new metric or experiment result is written to docs/experiments/ with the dataset name and model name.
- The README/docs are updated when a command or flow changes.

9 Experiment log template (docs/experiments/EXX_name.md)

```
# EXX: <question>
Date · dataset snapshot + sha · model(s) · code commit · Kaggle notebook version
## Hypothesis
## Setup (what changed vs baseline)
## Results (table: WER / CER / switch-point / halluc / RTF, by domain & flag)
## Error analysis –(1030 worst segments, categorised)
## Decision
```

10 Risks to the schedule

Risk	Early signal	Response
The aligner struggles on noisy or overlapping recordings	P1.7 precision < 70%	Lower θ , lean on VAD fallback for those recordings, prioritise clean recordings for v1
Review is slower than expected	P2.7 < 2× speed-up	Enable spot-check batch approval for segments with confidence > 0.9

Risk	Early signal	Response
Kaggle quota exhausted	Experiments queue up	Reduce to whisper-small only; shorter runs (2k steps); drop E5
v1 test WER far above 20%	P4.5	Error analysis first: label noise (fix data), translation (E5), hallucination (more non-speech samples), accent (more data). Do not jump to bigger models until the category is known
Gradio limits the reviewer UX	P2.4 feedback	Custom Gradio component for the waveform only

11 Immediate next steps

1. You: confirm the decisions flagged in the reply (interviewer turns, consent scope for the research audio, style-guide conventions).
2. Start Phase 0 on this branch: P0.1–P0.5 first (package, legacy move, the two bug fixes, CI).
3. You: pick the ~10% of recordings for test_v1, spread across noise levels and code-switch density, and keep them aside from here on.