

Transcription Style Guide

Status	v0.2: decisions confirmed; examples to be extended in Phase 0 (P0.8)
Version	1.2 · 2026-09-24 22:26 · change log
Applies to	Source transcripts (DOCX), reviewer edits in Corpus Studio, and training targets

Contents

1	Two layers of text	1
2	Language marking in DOCX	1
3	Verbatim rules	2
4	Numbers	2
5	Other spelling conventions	3
	Change log	4

1 Two layers of text

Layer	Where	Numbers	Purpose
Source text	The DOCX transcripts	Digits (6, 2019, Level 5)	Easy for people to write and read. Never modified by the pipeline
Spoken form	Segment revisions → training targets	Words in the language actually spoken (sita or six, twenty nineteen, Level Five)	What the model hears and learns
Display form	Transcribe output (optional)	Digits again, via inverse text normalisation	Readable subtitles and documents

The pipeline converts source text to spoken form automatically (§4). You keep writing digits in the documents.

2 Language marking in DOCX

- **Italic = Kiswahili**; roman (upright) = English. This applies to every document.
- **Hybrid words**: write them as **one word, no hyphen**. Italicise only the Kiswahili part if you can: *nimedownload*, *wameorganize*, *tutatextiana*. The parser tags a word that is part italic and part roman as mixed. If splitting the formatting is too slow, italicising the whole hybrid word is acceptable; the reviewer can re-tag it.
- Sheng words written with Kiswahili grammar are italic. English loanwords already absorbed into Kiswahili spelling (*kompyuta*, *basi*) are italic.
- Words that are neither English nor Kiswahili (e.g. Dholuo): write them as heard and wrap them in [luo: ...], e.g. [luo: ber ahinya].

3 Verbatim rules

Changed in v1.2: see the change log.

- Transcribe what was said, including repetitions that carry meaning and every code-switch.
- Drop pure fillers (*eh*, *umm*, *aah*) unless they carry meaning (*eeh* = “yes”).
- Unintelligible speech: [Unclear 13:10-13:15], with the start and end time (h:mm:ss after the first hour: [Unclear 1:05:21-1:05:24]); add a note if useful: [Unclear - Luo, 37:39-38:45]. Do not guess. These segments are excluded from training until resolved.
- Long stretches without transcribable speech: [No speech - activity, 17:26-18:14].
- **Always include the times** when a marked stretch is longer than about 2 seconds. The aligner uses them as fixed points, so words cannot drift into the stretch, and the report uses them to measure alignment accuracy. Use square brackets only: the tools also accept (Unclear: 1:16:07-1:16:08]], but mixed brackets are easy to get wrong.
- Participant references (P7, P12) stand in for a participant’s name for anonymity. They are *not* read as “pee seven”. Chunks containing one are kept out of training, because the audio holds a name the text does not.
- False starts and cut-off words: tuta--tutaanza, tuna--. The double hyphen marks the break and the parts are treated as separate words.
- Crosstalk: [overlap] at the start of the overlapping part.
- Passages you paraphrased or skipped for meaning: [nv] ... [/nv] (non-verbatim). These segments are excluded from training until corrected in review.

4 Numbers

Source documents always use **digits**. The spoken form uses **words in the language actually spoken**.

4.1 Default language rules (applied automatically)

Case	Rule	Source → spoken
Number inside an <i>italic</i> (Kiswahili) span	Kiswahili words	<i>watu 6</i> → watu sita
Number inside a roman (English) span	English words	6 people → six people
Year (4 digits, 1900–2099)	English words, year style	2019 → twenty nineteen; 2005 → two thousand and five
Label + number (Level, Grade, Form, Class, Chapter, Room, Route, Phase...)	English words, title case	Level 5 → Level Five; Form 4 → Form Four
Money	Follow the span language; the currency word stays as spoken	<i>shilingi 500</i> → shilingi mia tano; 500 bob → five hundred bob
Phone and ID numbers	Digit by digit, span language	0712 → zero seven one two
Time	Follow the span language	<i>saa 3</i> → saa tatu; 3 pm → three pm

4.2 Why a rule alone isn't enough, and how it gets resolved

A bare 6 can't tell you whether the speaker said *sita* or *six*. Kenyans often say numbers, and especially years and money, in English inside Kiswahili sentences. So the pipeline doesn't trust the rule blindly:

1. It generates both candidates (*sita* and *six*; twenty nineteen and elfu mbili kumi na tisa).
2. The forced aligner scores each candidate against the audio, and the higher-scoring one wins.
3. If the two scores are too close, the segment gets a `number_check` flag and the reviewer hears it and picks.

In practice the audio decides, and the §4.1 table is only the tie-breaker.

4.3 Kiswahili number words (reference)

1 moja · 2 mbili · 3 tatu · 4 nne · 5 tano · 6 sita · 7 saba · 8 nane · 9 tisa · 10 kumi · 11 kumi na moja · 20 ishirini · 30 thelathini · 40 arobaini · 50 hamsini · 60 sitini · 70 sabini · 80 themanini · 90 tisini · 100 mia moja · 250 mia mbili hamsini · 1,000 elfu moja · 2019 elfu mbili kumi na tisa · 1,000,000 milioni moja

Noun-class agreement (*watu watatu*, *vitu vitatu*) is written **as spoken** by the reviewer. The automatic converter emits the bare counting form (*tatu*), and the aligner or reviewer corrects it to the agreeing form.

5 Other spelling conventions

Item	Convention
Casing	Sentence case; proper nouns capitalised in both languages
Punctuation	Normal punctuation in transcripts (Whisper learns it); ignored when scoring
Contractions	As spoken: <i>don't</i> , <i>I'm</i> ; Kiswahili elisions as spoken (<i>nshaenda</i>)
Okay	<i>okay</i> (not <i>ok</i> , <i>OK</i> , <i>sawa</i> unless <i>sawa</i> was said)
Acronyms	Uppercase, no dots: <i>KCSE</i> , <i>NHIF</i> ; letter-by-letter acronyms stay uppercase

Change log

Times are East Africa Time (EAT, UTC+3). Each entry names the affected section, quotes what it said before and summarises what it says now, and links to the previous text: on GitHub at the exact commit and lines, and as an archived PDF of the previous version.

Version	Date (EAT)	Commit	Summary
1.0	2026-09-24 00:36	053e09d	First LaTeX version (v0.1 of the guide).
1.2	2026-09-24 22:26	c28ea1d	Marker format with timestamps, participant references, false starts (from the NRCCW_KSM09 pilot).

v1.2 · §3 Verbatim rules

Previous text: [v1.1 source, lines 43–57](#) · [v1.1 PDF](#)

Before. “Unintelligible speech: [unclear]. Do not guess.”; nothing on participant references, no-speech stretches or false starts.

Now. Markers carry start and end times ([Unclear 13:10-13:15], [No speech - activity, 17:26-18:14]), which the aligner uses as fixed points; square brackets only. P7-style participant references are anonymised names, not spoken letters, and keep their chunk out of training. False starts use --.